

Zhuowei Chen

◇ Email: zhuowei.chen@pitt.edu | ◇ +1 (412)-589-7039

<https://johnnychanv.github.io>

EDUCATION

University of Pittsburgh Ph.D. in Computer Science Advisor: Prof. Xiang Lorraine Li	Pittsburgh, PA Sep 2025 - Current
Guangdong University of Foreign Studies (GPA: 3.82) Bachelor in Software Engineering Advisor: Prof. Lianxi Wang	Guangzhou, China Sep 2021 - Jun 2025
University of California, Berkeley (GPA: 4.00) Visiting Student	Berkeley, CA Aug 2023 - Jun 2024

RESEARCH EXPERIENCE

Oracle America, Inc. - Oracle AI <i>Applied Scientist Intern</i> Supervisor: Dr. Avi Sil & Dr. Yingshen Wang	Seattle, WA Sep 2026 – Current
- Built an agentic benchmark for self-evolving operations-research (OR) agents in real industrial settings. - Diagnosed agent failure modes via neuron-level analysis of trajectories from eight LLMs.	
University of Pittsburgh <i>Graduate Student Researcher</i> Supervisor: Prof. Xiang Lorraine Li	Pittsburgh, PA Sep 2025 – Current
- Selected few-shot samples for annotation from neuron activation patterns rather than output-level entropy. - Built an annotation-free self-distillation pipeline with neuron-guided data selection and teacher context.	
Wisers Information Ltd. - Wisers AI Lab <i>NLP Research Intern</i>	Shenzhen, China Dec 2023 - Mar 2024
- Built Roberta-based textual classification models with human-annotated social media content. - Applied transformers for time series regression to predict regional arrivals.	

SELECTED PUBLICATIONS

* represents equal contributions and † represents the corresponding author.

Conference Papers

- Zhuowei Chen**, Liwei Chen, Christian Schunn, Raquel Coelho, Xiang Lorraine Li.
[Neuron-Aware Active Few-Shot Learning for LLMs.](#)
The 64th Annual Meeting of the Association for Computational Linguistics, ACL 2026 (Oral).
Data Selection LLM Adaptation
- Zhuowei Chen**, Qiannan Zhang, Shichao Pei.
[Injecting Universal Jailbreak Backdoors to LLMs in Minutes.](#)
The Thirteenth International Conference on Learning Representations, ICLR 2025.
Model Editing LLM Safety
- Zhuowei Chen**, Yuben Wu, Xinfeng Liao, Yujia Tian, Lianxi Wang[†].
[An Effective Deployment of Diffusion LM for Data Augmentation in Low-Resource Sentiment Classification.](#)
The 2024 Conference on Empirical Methods in Natural Language Processing, EMNLP 2024 (Main).
Data Augmentation Diffusion LM

4.

Lianxi Wang, Yujia Tian*, **Zhuowei Chen***[†].
Enhancing Hindi Feature Representation Through Fusion of Dual-Script Word Embeddings.
The 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation, LREC-COLING 2024 (Main).

MoE RepresentationLow-Resource NLP

5.

Zhuowei Chen, Yujia Tian, Lianxi Wang[†], Shengyi Jiang.
A Distantly-Supervised Relation Extraction Method Based on Selective Gate and Noise Correction.
The 22nd China National Conference on Computational Linguistics, CCL 2023.

Feature SelectionLow-Resource NLP

Workshop and Preprints

1.

Zhaoyi Joey Hou, **Zhuowei Chen**, Mengxue Zhang, Xiang Lorraine Li.
Rethinking Post-training Diversity Collapse: Is Diversity-preserving Post-training Enough?
Human-AI Co-Creativity Workshop @ ICML 2026.

LLM Post-trainingLLM Creativity

2.

Zhuowei Chen, Xiang Lorraine Li.
Neuron-Aware Data Selection for Annotation-Free LLM Self-Distillation.
DEMO Workshop @ ICML 2026.

Data SelectionLLM Adaptation

3.

Zhuowei Chen, Bowei Zhang, Nankai Lin, Tian Hou, Lianxi Wang.
LLM Safeguards for Low-Resource Languages via Reasoning and Alignment with Minimal Training Data.
Multilingual Representation Learning Workshop @ EMNLP 2025.

LLM Post-trainingLLM Safety

SELECTED HONORS

- Top Ten Outstanding Youth

(Top 10/30000) Guangdong University of Foreign Studies, 2025
- China National Scholarship

(Top 0.2%) Ministry of Education of the PRC, 2024

SERVICES

Reviewer

- IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)

2026
- The International Conference on Learning Representations (ICLR)

2025
- ACL Rolling Review (ARR)

2024, 2025, 2026